



# **Deleting failed plexes owned by the surviving site (MetroCluster IP configurations)**

## **ONTAP MetroCluster**

Paula Carrigan, Thom Illingworth  
June 21, 2021

This PDF was generated from [https://docs.netapp.com/us-en/ontap-metrocluster/disaster-recovery/task\\_delete\\_plexes\\_mcip.html](https://docs.netapp.com/us-en/ontap-metrocluster/disaster-recovery/task_delete_plexes_mcip.html) on September 24, 2021. Always check docs.netapp.com for the latest.

# Table of Contents

Deleting failed plexes owned by the surviving site (MetroCluster IP configurations) . . . . . 1

# Deleting failed plexes owned by the surviving site (MetroCluster IP configurations)

After replacing hardware and assigning disks, you must delete failed remote plexes that are owned by the surviving site nodes but located at the disaster site.

## About this task

These steps are performed on the surviving cluster.

## Steps

1. Identify the local aggregates: `storage aggregate show -is-home true`

```
cluster_B::> storage aggregate show -is-home true

cluster_B Aggregates:
Aggregate      Size Available Used% State   #Vols  Nodes           RAID
Status
-----
node_B_1_aggr0 1.49TB   74.12GB 95% online      1 node_B_1
raid4,

mirror

degraded
node_B_2_aggr0 1.49TB   74.12GB 95% online      1 node_B_2
raid4,

mirror

degraded
node_B_1_aggr1 2.99TB   2.88TB   3% online     15 node_B_1
raid_dp,

mirror

degraded
node_B_1_aggr2 2.99TB   2.91TB   3% online     14 node_B_1
raid_tec,

mirror

degraded
node_B_2_aggr1 2.95TB   2.80TB   5% online     37 node_B_2
raid_dp,
```

```
mirror

degraded
node_B_2_aggr2 2.99TB  2.87TB    4% online      35 node_B_2
raid_tec,

mirror

degraded
6 entries were displayed.

cluster_B::>
```

2. Identify the failed remote plexes:

```
storage aggregate plex show
```

The following example calls out the plexes that are remote (not plex0) and have a status of "failed":

```
cluster_B::> storage aggregate plex show -fields aggregate,status,is-
online,Plex,pool
aggregate      plex  status          is-online pool
-----
node_B_1_aggr0 plex0 normal,active true      0
node_B_1_aggr0 plex4 failed,inactive false - <<<<---Plex at remote site
node_B_2_aggr0 plex0 normal,active true      0
node_B_2_aggr0 plex4 failed,inactive false - <<<<---Plex at remote site
node_B_1_aggr1 plex0 normal,active true      0
node_B_1_aggr1 plex4 failed,inactive false - <<<<---Plex at remote site
node_B_1_aggr2 plex0 normal,active true      0
node_B_1_aggr2 plex1 failed,inactive false - <<<<---Plex at remote site
node_B_2_aggr1 plex0 normal,active true      0
node_B_2_aggr1 plex4 failed,inactive false - <<<<---Plex at remote site
node_B_2_aggr2 plex0 normal,active true      0
node_B_2_aggr2 plex1 failed,inactive false - <<<<---Plex at remote site
node_A_1_aggr1 plex0 failed,inactive false -
node_A_1_aggr1 plex4 normal,active true      1
node_A_1_aggr2 plex0 failed,inactive false -
node_A_1_aggr2 plex1 normal,active true      1
node_A_2_aggr1 plex0 failed,inactive false -
node_A_2_aggr1 plex4 normal,active true      1
node_A_2_aggr2 plex0 failed,inactive false -
node_A_2_aggr2 plex1 normal,active true      1
20 entries were displayed.

cluster_B::>
```

### 3. Take offline each of the failed plexes, and then delete them:

#### a. Take offline the failed plexes:

```
storage aggregate plex offline -aggregate aggregate-name -plex plex-id
```

The following example shows the aggregate "node\_B\_2\_aggr1/plex1" being taken offline:

```
cluster_B::> storage aggregate plex offline -aggregate node_B_1_aggr0
-plex plex4

Plex offline successful on plex: node_B_1_aggr0/plex4
```

#### b. Delete the failed plex:

```
storage aggregate plex delete -aggregate aggregate-name -plex plex-id
```

You can destroy the plex when prompted.

The following example shows the plex node\_B\_2\_aggr1/plex1 being deleted.

```
cluster_B::> storage aggregate plex delete -aggregate node_B_1_aggr0
-plex plex4

Warning: Aggregate "node_B_1_aggr0" is being used for the local
management root
        volume or HA partner management root volume, or has been
marked as
        the aggregate to be used for the management root volume
after a
        reboot operation. Deleting plex "plex4" for this aggregate
could lead
        to unavailability of the root volume after a disaster
recovery
        procedure. Use the "storage aggregate show -fields
        has-mroot,has-partner-mroot,root" command to view such
aggregates.

Warning: Deleting plex "plex4" of mirrored aggregate "node_B_1_aggr0"
on node
        "node_B_1" in a MetroCluster configuration will disable its
synchronous disaster recovery protection. Are you sure you
want to
        destroy this plex? {y|n}: y
[Job 633] Job succeeded: DONE

cluster_B::>
```

You must repeat these steps for each of the failed plexes.

#### 4. Confirm that the plexes have been removed:

```
storage aggregate plex show -fields aggregate,status,is-online,plex,pool
```

```
cluster_B::> storage aggregate plex show -fields aggregate,status,is-
online,Plex,pool
aggregate      plex  status          is-online pool
-----
node_B_1_aggr0 plex0 normal,active true      0
node_B_2_aggr0 plex0 normal,active true      0
node_B_1_aggr1 plex0 normal,active true      0
node_B_1_aggr2 plex0 normal,active true      0
node_B_2_aggr1 plex0 normal,active true      0
node_B_2_aggr2 plex0 normal,active true      0
node_A_1_aggr1 plex0 failed,inactive false    -
node_A_1_aggr1 plex4 normal,active true      1
node_A_1_aggr2 plex0 failed,inactive false    -
node_A_1_aggr2 plex1 normal,active true      1
node_A_2_aggr1 plex0 failed,inactive false    -
node_A_2_aggr1 plex4 normal,active true      1
node_A_2_aggr2 plex0 failed,inactive false    -
node_A_2_aggr2 plex1 normal,active true      1
14 entries were displayed.

cluster_B::>
```

##### 5. Identify the switched-over aggregates:

```
storage aggregate show -is-home false
```

You can also use the `storage aggregate plex show -fields aggregate,status,is-online,plex,pool` command to identify plex 0 switched-over aggregates. They will have a status of "failed, inactive".

The following commands show four switched-over aggregates:

- `node_A_1_aggr1`
- `node_A_1_aggr2`
- `node_A_2_aggr1`
- `node_A_2_aggr2`

```

cluster_B::> storage aggregate show -is-home false

cluster_A Switched Over Aggregates:
Aggregate      Size Available Used% State   #Vols  Nodes      RAID
Status
-----
node_A_1_aggr1 2.12TB  1.88TB   11% online    91 node_B_1
raid_dp,

mirror

degraded
node_A_1_aggr2 2.89TB  2.64TB    9% online    90 node_B_1
raid_tec,

mirror

degraded
node_A_2_aggr1 2.12TB  1.86TB   12% online    91 node_B_2
raid_dp,

mirror

degraded
node_A_2_aggr2 2.89TB  2.64TB    9% online    90 node_B_2
raid_tec,

mirror

degraded
4 entries were displayed.

cluster_B::>

```

#### 6. Identify switched-over plexes:

```
storage aggregate plex show -fields aggregate,status,is-online,Plex,pool
```

You want to identify the plexes with a status of "failed, inactive".

The following commands show four switched-over aggregates:



```
cluster_B::> storage aggregate plex show -fields aggregate,status,is-
online,Plex,pool
aggregate      plex  status          is-online pool
-----
node_B_1_aggr0 plex0 normal,active true      0
node_B_2_aggr0 plex0 normal,active true      0
node_B_1_aggr1 plex0 normal,active true      0
node_B_1_aggr2 plex0 normal,active true      0
node_B_2_aggr1 plex0 normal,active true      0
node_B_2_aggr2 plex0 normal,active true      0
node_A_1_aggr1 plex0 failed,inactive false - <<<<-- Switched over
aggr/Plex0
node_A_1_aggr1 plex4 normal,active true      1
node_A_1_aggr2 plex0 failed,inactive false - <<<<-- Switched over
aggr/Plex0
node_A_1_aggr2 plex1 normal,active true      1
node_A_2_aggr1 plex0 failed,inactive false - <<<<-- Switched over
aggr/Plex0
node_A_2_aggr1 plex4 normal,active true      1
node_A_2_aggr2 plex0 failed,inactive false - <<<<-- Switched over
aggr/Plex0
node_A_2_aggr2 plex1 normal,active true      1
14 entries were displayed.

cluster_B::>
```

## 7. Delete the failed plex:

```
storage aggregate plex delete -aggregate node_A_1_aggr1 -plex plex0
```

You can destroy the plex when prompted.

The following example shows the plex node\_A\_1\_aggr1/plex0 being deleted:

```

cluster_B::> storage aggregate plex delete -aggregate node_A_1_aggr1
-plex plex0

Warning: Aggregate "node_A_1_aggr1" hosts MetroCluster metadata volume
"MDV_CRS_e8457659b8a711e78b3b00a0988fe74b_A". Deleting plex
"plex0"
      for this aggregate can lead to the failure of configuration
      replication across the two DR sites. Use the "volume show
-vserver
      <admin-vserver> -volume MDV_CRS*" command to verify the
location of
      such volumes.

Warning: Deleting plex "plex0" of mirrored aggregate "node_A_1_aggr1" on
node
      "node_A_1" in a MetroCluster configuration will disable its
      synchronous disaster recovery protection. Are you sure you want
to
      destroy this plex? {y|n}: y
[Job 639] Job succeeded: DONE

cluster_B::>

```

You must repeat these steps for each of the failed aggregates.

#### 8. Verify that there are no failed plexes remaining on the surviving site.

The following output shows that all plexes are normal, active, and online.

```
cluster_B::> storage aggregate plex show -fields aggregate,status,is-  
online,Plex,pool
```

aggregate	plex	status	is-online	pool
-----------	------	--------	-----------	------

-----	-----	-----	-----	-----
-------	-------	-------	-------	-------

node_B_1_aggr0	plex0	normal,active	true	0
node_B_2_aggr0	plex0	normal,active	true	0
node_B_1_aggr1	plex0	normal,active	true	0
node_B_2_aggr2	plex0	normal,active	true	0
node_B_1_aggr1	plex0	normal,active	true	0
node_B_2_aggr2	plex0	normal,active	true	0
node_A_1_aggr1	plex4	normal,active	true	1
node_A_1_aggr2	plex1	normal,active	true	1
node_A_2_aggr1	plex4	normal,active	true	1
node_A_2_aggr2	plex1	normal,active	true	1

10 entries were displayed.

```
cluster_B::>
```

## Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system- without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

## Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.