



Configure the MetroCluster for transition

ONTAP MetroCluster

Martin Houser, Ranu Kundu
June 21, 2021

This PDF was generated from https://docs.netapp.com/us-en/ontap-metrocluster/transition/task_configure_transition.html on September 24, 2021. Always check docs.netapp.com for the latest.

Table of Contents

- Configure the MetroCluster for transition 1
 - Sending a custom AutoSupport message prior to maintenance 1
 - Enabling transition mode and disabling cluster HA 1
 - Joining the MetroCluster IP nodes to the clusters 4
 - Configuring intercluster LIFs, creating the MetroCluster interfaces, and mirroring root aggregates 6
 - Finalizing the addition of the MetroCluster IP nodes 16

Configure the MetroCluster for transition

To prepare the configuration for transition you add the new nodes to the existing MetroCluster configuration and then move data to the new nodes.

Sending a custom AutoSupport message prior to maintenance

Before performing the maintenance, you should issue an AutoSupport message to notify NetApp technical support that maintenance is underway. Informing technical support that maintenance is underway prevents them from opening a case on the assumption that a disruption has occurred.

This task must be performed on each MetroCluster site.

1. To prevent automatic support case generation, send an Autosupport message to indicate maintenance is underway.

- a. Issue the following command: `system node autosupport invoke -node * -type all -message MAINT=maintenance-window-in-hours`

`maintenance-window-in-hours` specifies the length of the maintenance window, with a maximum of 72 hours. If the maintenance is completed before the time has elapsed, you can invoke an AutoSupport message indicating the end of the maintenance period: `system node autosupport invoke -node * -type all -message MAINT=end`

- b. Repeat the command on the partner cluster.

Enabling transition mode and disabling cluster HA

You must enable the MetroCluster transition mode to allow the old and new nodes to operate together in the MetroCluster configuration, and disable cluster HA.

1. Enable transition:
 - a. Change to the advanced privilege level: `set -privilege advanced`
 - b. Enable transition mode: `metrocluster transition enable -transition-mode non-disruptive`



Run this command on one cluster only.

```
cluster_A::*> metrocluster transition enable -transition-mode non-disruptive
```

Warning: This command enables the start of a "non-disruptive" MetroCluster

FC-to-IP transition. It allows the addition of hardware for another DR

group that uses IP fabrics, and the removal of a DR group that uses FC

fabrics. Clients will continue to access their data during a non-disruptive transition.

Automatic unplanned switchover will also be disabled by this command.

Do you want to continue? {y|n}: y

```
cluster_A::*>
```

c. Return to the admin privilege level: `set -privilege admin`

2. Verify that transition is enabled on both the clusters.

```
cluster_A::> metrocluster transition show-mode  
Transition Mode
```

non-disruptive

```
cluster_A::*>
```

```
cluster_B::*> metrocluster transition show-mode  
Transition Mode
```

non-disruptive

```
Cluster_B::>
```

3. Disable cluster HA.



You must run this command on both clusters.

```
cluster_A::~*> cluster ha modify -configured false
```

Warning: This operation will unconfigure cluster HA. Cluster HA must be configured on a two-node cluster to ensure data access availability in the event of storage failover.

Do you want to continue? {y|n}: y

Notice: HA is disabled.

```
cluster_A::~*>
```

```
cluster_B::~*> cluster ha modify -configured false
```

Warning: This operation will unconfigure cluster HA. Cluster HA must be configured on a two-node cluster to ensure data access availability in the event of storage failover.

Do you want to continue? {y|n}: y

Notice: HA is disabled.

```
cluster_B::~*>
```

4. Verify that cluster HA is disabled.



You must run this command on both clusters.

```
cluster_A:> cluster ha show
```

```
High Availability Configured: false
```

```
Warning: Cluster HA has not been configured. Cluster HA must be  
configured
```

```
on a two-node cluster to ensure data access availability in the  
event of storage failover. Use the "cluster ha modify -configured  
true" command to configure cluster HA.
```

```
cluster_A:>
```

```
cluster_B:> cluster ha show
```

```
High Availability Configured: false
```

```
Warning: Cluster HA has not been configured. Cluster HA must be  
configured
```

```
on a two-node cluster to ensure data access availability in the  
event of storage failover. Use the "cluster ha modify -configured  
true" command to configure cluster HA.
```

```
cluster_B:>
```

Joining the MetroCluster IP nodes to the clusters

You must add the four new MetroCluster IP nodes to the existing MetroCluster configuration.

You must perform this task on both clusters.

1. Add the MetroCluster IP nodes to the existing MetroCluster configuration.
 - a. Join the first MetroCluster IP node (node_A_1-IP) to the existing MetroCluster FC configuration.

```
Welcome to the cluster setup wizard.
```

```
You can enter the following commands at any time:
```

```
"help" or "?" - if you want to have a question clarified,  
"back" - if you want to change previously answered questions, and  
"exit" or "quit" - if you want to quit the cluster setup wizard.  
Any changes you made before quitting will be saved.
```

```
You can return to cluster setup at any time by typing "cluster  
setup".
```

```
To accept a default or omit a question, do not enter a value.
```

This system will send event messages and periodic reports to NetApp Technical Support. To disable this feature, enter `autosupport modify -support disable` within 24 hours.

Enabling AutoSupport can significantly speed problem determination and resolution, should a problem occur on your system. For further information on AutoSupport, see: <http://support.netapp.com/autosupport/>

Type yes to confirm and continue {yes}: yes

Enter the node management interface port [e0M]:
Enter the node management interface IP address: 172.17.8.93
Enter the node management interface netmask: 255.255.254.0
Enter the node management interface default gateway: 172.17.8.1
A node management interface on port e0M with IP address 172.17.8.93 has been created.

Use your web browser to complete cluster setup by accessing <https://172.17.8.93>

Otherwise, press Enter to complete cluster setup using the command line interface:

Do you want to create a new cluster or join an existing cluster? {create, join}:
join

Existing cluster interface configuration found:

Port	MTU	IP	Netmask
e0c	9000	169.254.148.217	255.255.0.0
e0d	9000	169.254.144.238	255.255.0.0

Do you want to use this configuration? {yes, no} [yes]: yes
.
.
.

- b. Join the second MetroCluster IP node (node_A_2-IP) to the existing MetroCluster FC configuration.
2. Repeat these steps to join node_B_1-IP and node_B_2-IP to cluster_B.

Configuring intercluster LIFs, creating the MetroCluster interfaces, and mirroring root aggregates

You must create cluster peering LIFs, create the MetroCluster interfaces on the new MetroCluster IP nodes.

The home port used in the examples are platform-specific. You should use the appropriate home port specific to MetroCluster IP node platform.

1. On the new MetroCluster IP nodes, [configure the intercluster LIFs](#).
2. On each site, verify that cluster peering is configured: `cluster peer show`

The following example shows the cluster peering configuration on cluster_A:

```
cluster_A:> cluster peer show
Peer Cluster Name      Cluster Serial Number Availability
Authentication
-----
cluster_B              1-80-000011          Available      ok
```

The following example shows the cluster peering configuration on cluster_B:

```
cluster_B:> cluster peer show
Peer Cluster Name      Cluster Serial Number Availability
Authentication
-----
cluster_A 1-80-000011  Available      ok
```

3. Configure the DR group for the MetroCluster IP nodes: `metrocluster configuration-settings dr-group create -partner-cluster`

```
cluster_A::> metrocluster configuration-settings dr-group create
-partner-cluster
cluster_B -local-node node_A_3-IP -remote-node node_B_3-IP
[Job 259] Job succeeded: DR Group Create is successful.
cluster_A::>
```

4. Verify that the DR group is created. `metrocluster configuration-settings dr-group show`


```

cluster_A::> metrocluster configuration-settings dr-group show

DR Group ID Cluster                               Node                               DR Partner
Node
-----
2          cluster_A
          node_A_3-IP                          node_B_3-IP
          node_A_4-IP                          node_B_4-IP
          cluster_B
          node_B_3-IP                          node_A_3-IP
          node_B_4-IP                          node_A_4-IP

4 entries were displayed.

cluster_A::>

```

You will notice that the DR group for the old MetroCluster FC nodes (DR Group 1) is not listed when you run the `metrocluster configuration-settings dr-group show` command.

You can use `metrocluster node show` command on both sites to list all nodes.

```
cluster_A::> metrocluster node show
```

DR			Configuration	DR	
Group	Cluster	Node	State	Mirroring	Mode
-----	-----	-----	-----	-----	-----
1	cluster_A				
		node_A_1-FC	configured	enabled	normal
		node_A_2-FC	configured	enabled	normal
	cluster_B				
		node_B_1-FC	configured	enabled	normal
		node_B_2-FC	configured	enabled	normal
2	cluster_A				
		node_A_1-IP	ready to configure	-	-
		node_A_2-IP	ready to configure	-	-

```
cluster_B::> metrocluster node show
```

DR			Configuration	DR	
Group	Cluster	Node	State	Mirroring	Mode
-----	-----	-----	-----	-----	-----
1	cluster_B				
		node_B_1-FC	configured	enabled	normal
		node_B_2-FC	configured	enabled	normal
	cluster_A				
		node_A_1-FC	configured	enabled	normal
		node_A_2-FC	configured	enabled	normal
2	cluster_B				
		node_B_1-IP	ready to configure	-	-
		node_B_2-IP	ready to configure	-	-

5. Configure the MetroCluster IP interfaces for the newly joined MetroCluster IP nodes: `metrocluster configuration-settings interface create -cluster-name`

See [Configuring and connecting the MetroCluster IP interfaces](#) for considerations when configuring the IP interfaces. NOTE: You can configure the MetroCluster IP interfaces from either cluster. Also, starting with ONTAP 9.9.1, if you are using a layer 3 configuration, you must also specify the `-gateway` parameter when creating MetroCluster IP interfaces. Refer to [xref:./transition/./install-ip/concept_considerations_layer_3.html](#).

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_3-IP -home-port elb -address
172.17.26.10 -netmask 255.255.255.0
[Job 260] Job succeeded: Interface Create is successful.
```

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_3-IP -home-port elb -address
172.17.27.10 -netmask 255.255.255.0
[Job 261] Job succeeded: Interface Create is successful.
```

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_4-IP -home-port elb -address
172.17.26.11 -netmask 255.255.255.0
[Job 262] Job succeeded: Interface Create is successful.
```

```
cluster_A::> :metrocluster configuration-settings interface create
-cluster-name cluster_A -home-node node_A_4-IP -home-port elb -address
172.17.27.11 -netmask 255.255.255.0
[Job 263] Job succeeded: Interface Create is successful.
```

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_3-IP -home-port elb -address
172.17.26.12 -netmask 255.255.255.0
[Job 264] Job succeeded: Interface Create is successful.
```

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_3-IP -home-port elb -address
172.17.27.12 -netmask 255.255.255.0
[Job 265] Job succeeded: Interface Create is successful.
```

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_4-IP -home-port elb -address
172.17.26.13 -netmask 255.255.255.0
[Job 266] Job succeeded: Interface Create is successful.
```

```
cluster_A::> metrocluster configuration-settings interface create
-cluster-name cluster_B -home-node node_B_4-IP -home-port elb -address
172.17.27.13 -netmask 255.255.255.0
[Job 267] Job succeeded: Interface Create is successful.
```

6. Verify the MetroCluster IP interfaces are created: metrocluster configuration-settings interface show

```

cluster_A::>metrocluster configuration-settings interface show

DR
Config
Group Cluster Node      Network Address Netmask      Gateway
State
-----
-----
2      cluster_A
      node_A_3-IP
      Home Port: ela
      172.17.26.10      255.255.255.0      -
completed
      Home Port: elb
      172.17.27.10      255.255.255.0      -
completed
      node_A_4-IP
      Home Port: ela
      172.17.26.11      255.255.255.0      -
completed
      Home Port: elb
      172.17.27.11      255.255.255.0      -
completed
      cluster_B
      node_B_3-IP
      Home Port: ela
      172.17.26.13      255.255.255.0      -
completed
      Home Port: elb
      172.17.27.13      255.255.255.0      -
completed
      node_B_3-IP
      Home Port: ela
      172.17.26.12      255.255.255.0      -
completed
      Home Port: elb
      172.17.27.12      255.255.255.0      -
completed
8 entries were displayed.

cluster_A>

```

7. Connect the MetroCluster IP interfaces: metrocluster configuration-settings connection connect



This command might take several minutes to complete.

```
cluster_A::> metrocluster configuration-settings connection connect

cluster_A::>
```

8. Verify the connections are properly established: metrocluster configuration-settings connection show

```
cluster_A::> metrocluster configuration-settings connection show
```

DR Group	Cluster	Node	Source Network Address	Destination Network Address	Partner Type
2	cluster_A	node_A_3-IP**	Home Port: ela 172.17.26.10	172.17.26.11	HA Partner
completed			Home Port: ela 172.17.26.10	172.17.26.12	DR Partner
completed			Home Port: ela 172.17.26.10	172.17.26.13	DR Auxiliary
completed			Home Port: elb 172.17.27.10	172.17.27.11	HA Partner
completed			Home Port: elb 172.17.27.10	172.17.27.12	DR Partner
completed			Home Port: elb 172.17.27.10	172.17.27.13	DR Auxiliary
		node_A_4-IP	Home Port: ela 172.17.26.11	172.17.26.10	HA Partner
completed			Home Port: ela 172.17.26.11	172.17.26.13	DR Partner
completed			Home Port: ela		

```

completed          172.17.26.11      172.17.26.12      DR Auxiliary
                    Home Port: elb
                    172.17.27.11      172.17.27.10      HA Partner
completed
                    Home Port: elb
                    172.17.27.11      172.17.27.13      DR Partner
completed
                    Home Port: elb
                    172.17.27.11      172.17.27.12      DR Auxiliary
completed

DR                  Source          Destination
Group Cluster Node   Network Address Network Address Partner Type
Config State
-----
2      cluster_B
      node_B_4-IP
      Home Port: ela
      172.17.26.13      172.17.26.12      HA Partner
completed
      Home Port: ela
      172.17.26.13      172.17.26.11      DR Partner
completed
      Home Port: ela
      172.17.26.13      172.17.26.10      DR Auxiliary
completed
      Home Port: elb
      172.17.27.13      172.17.27.12      HA Partner
completed
      Home Port: elb
      172.17.27.13      172.17.27.11      DR Partner
completed
      Home Port: elb
      172.17.27.13      172.17.27.10      DR Auxiliary
completed
      node_B_3-IP
      Home Port: ela
      172.17.26.12      172.17.26.13      HA Partner
completed
      Home Port: ela
      172.17.26.12      172.17.26.10      DR Partner
completed
      Home Port: ela
      172.17.26.12      172.17.26.11      DR Auxiliary

```

```
completed
      Home Port: elb
      172.17.27.12    172.17.27.13    HA Partner
completed
      Home Port: elb
      172.17.27.12    172.17.27.10    DR Partner
completed
      Home Port: elb
      172.17.27.12    172.17.27.11    DR Auxiliary
completed
24 entries were displayed.

cluster_A::>
```

9. Verify disk autoassignment and partitioning: `disk show -pool Pool1`

```
cluster_A::> disk show -pool Pool1
```

Disk Owner	Usable Size	Shelf	Bay	Disk Type	Container Type	Container Name
-----	-----	-----	---	-----	-----	-----
1.10.4	-	10	4	SAS	remote	-
node_B_2						
1.10.13	-	10	13	SAS	remote	-
node_B_2						
1.10.14	-	10	14	SAS	remote	-
node_B_1						
1.10.15	-	10	15	SAS	remote	-
node_B_1						
1.10.16	-	10	16	SAS	remote	-
node_B_1						
1.10.18	-	10	18	SAS	remote	-
node_B_2						
...						
2.20.0	546.9GB	20	0	SAS	aggregate	aggr0_rha1_a1
node_a_1						
2.20.3	546.9GB	20	3	SAS	aggregate	aggr0_rha1_a2
node_a_2						
2.20.5	546.9GB	20	5	SAS	aggregate	rha1_a1_aggr1
node_a_1						
2.20.6	546.9GB	20	6	SAS	aggregate	rha1_a1_aggr1
node_a_1						
2.20.7	546.9GB	20	7	SAS	aggregate	rha1_a2_aggr1
node_a_2						
2.20.10	546.9GB	20	10	SAS	aggregate	rha1_a1_aggr1
node_a_1						
...						

43 entries were displayed.

```
cluster_A::>
```

10. Mirror the root aggregates: `storage aggregate mirror -aggregate aggr0_node_A_3-IP`



You must complete this step on each MetroCluster IP node.


```
cluster_A::> aggr mirror -aggregate aggr0_node_A_3-IP
```

Info: Disks would be added to aggregate "aggr0_node_A_3-IP" on node "node_A_3-IP" in the following manner:

Second Plex

RAID Group rg0, 3 disks (block checksum, raid_dp)

Physical Size	Position	Disk	Type	Usable Size
-----	-----	-----	-----	-----
-----	dparity	4.20.0	SAS	-
-	parity	4.20.3	SAS	-
-	data	4.20.1	SAS	546.9GB
558.9GB				

Aggregate capacity available for volume use would be 467.6GB.

Do you want to continue? {y|n}: y

```
cluster_A::>
```

11. Verify that the root aggregates are mirrored: `storage aggregate show`

```
cluster_A::> aggr show
```

Aggregate Status	Size	Available	Used%	State	#Vols	Nodes	RAID
-----	-----	-----	-----	-----	-----	-----	-----

aggr0_node_A_1-FC	349.0GB	16.84GB	95%	online	1	node_A_1-FC	
raid_dp,							
mirrored,							
normal							
aggr0_node_A_2-FC	349.0GB	16.84GB	95%	online	1	node_A_2-FC	

```

raid_dp,

mirrored,

normal
aggr0_node_A_3-IP
          467.6GB    22.63GB    95% online          1 node_A_3-IP
raid_dp,

mirrored,

normal
aggr0_node_A_4-IP
          467.6GB    22.62GB    95% online          1 node_A_4-IP
raid_dp,

mirrored,

normal
aggr_data_a1
          1.02TB     1.01TB     1% online          1 node_A_1-FC
raid_dp,

mirrored,

normal
aggr_data_a2
          1.02TB     1.01TB     1% online          1 node_A_2-FC
raid_dp,

mirrored,

```

Finalizing the addition of the MetroCluster IP nodes

You must incorporate the new DR group into the MetroCluster configuration and create mirrored data aggregates on the new nodes.

1. Create mirrored data aggregates on each of the new MetroCluster nodes: `storage aggregate create -aggregate aggregate-name -node node-name -diskcount no-of-disks -mirror true`



You must create at least one mirrored data aggregate per site. It is recommended to have two mirrored data aggregates per site on MetroCluster IP nodes to host the MDV volumes, however a single aggregate per site is supported (but not recommended). It is support that one site of the MetroCluster has a single mirrored data aggregate and the other site has more than one mirrored data aggregate.

The following example shows the creation of an aggregate on node_A_1-new.

```
cluster_A::> storage aggregate create -aggregate data_a3 -node node_A_1-  
new -diskcount 10 -mirror t
```

Info: The layout for aggregate "data_a3" on node "node_A_1-new" would be:

First Plex

RAID Group rg0, 5 disks (block checksum, raid_dp)

				Usable
Physical	Position	Disk	Type	Size
Size				
-----	-----	-----	-----	-----
-----	dparity	5.10.15	SAS	-
-	parity	5.10.16	SAS	-
-	data	5.10.17	SAS	546.9GB
547.1GB	data	5.10.18	SAS	546.9GB
558.9GB	data	5.10.19	SAS	546.9GB
558.9GB				

Second Plex

RAID Group rg0, 5 disks (block checksum, raid_dp)

				Usable
Physical	Position	Disk	Type	Size
Size				
-----	-----	-----	-----	-----
-----	dparity	4.20.17	SAS	-
-	parity	4.20.14	SAS	-
-	data	4.20.18	SAS	546.9GB
547.1GB	data	4.20.19	SAS	546.9GB
547.1GB	data	4.20.16	SAS	546.9GB

```
547.1GB
```

```
Aggregate capacity available for volume use would be 1.37TB.
```

```
Do you want to continue? {y|n}: y
```

```
[Job 440] Job succeeded: DONE
```

```
cluster_A::>
```

2. Configure the MetroCluster to implement the changes: `metrocluster configure`

```
cluster_A::*> metrocluster configure
```

```
[Job 439] Job succeeded: Configure is successful.
```

```
cluster_A::*>
```

3. Verify that the nodes are added to their DR group: `metrocluster node show`

```
cluster_A::*> metrocluster node show
```

DR Group	Cluster	Node	Configuration State	DR Mirroring Mode

1	cluster_A			
		node-A-1-FC	configured	enabled normal
		node-A-2-FC	configured	enabled normal
	Cluster-B			
		node-B-1-FC	configured	enabled normal
		node-B-2-FC	configured	enabled normal
2	cluster_A			
		node-A-3-IP	configured	enabled normal
		node-A-4-IP	configured	enabled normal
	Cluster-B			
		node-B-3-IP	configured	enabled normal
		node-B-4-IP	configured	enabled normal

8 entries were displayed.

```
cluster_A::~*>
```

4. Move the MDV_CRS volumes from the old nodes to the new nodes in advanced privilege.

a. Display the volumes to identify the MDV volumes:



If you have a single mirrored data aggregate per site then move both the MDV volumes to this single aggregate. If you have two or more mirrored data aggregates, then move each MDV volume to a different aggregate.

The following example shows the MDV volumes in the volume show output:

```
cluster_A::> volume show
Vserver  Volume                Aggregate  State  Type  Size
Available Used%
-----
...

cluster_A  MDV_CRS_2c78e009ff5611e9b0f300a0985ef8c4_A
          aggr_b1          -        RW        -
-      -
cluster_A  MDV_CRS_2c78e009ff5611e9b0f300a0985ef8c4_B
          aggr_b2          -        RW        -
-      -
cluster_A  MDV_CRS_d6b0b313ff5611e9837100a098544e51_A
          aggr_a1        online    RW        10GB
9.50GB    0%
cluster_A  MDV_CRS_d6b0b313ff5611e9837100a098544e51_B
          aggr_a2        online    RW        10GB
9.50GB    0%
...
11 entries were displayed.mple
```

- b. Set the advanced privilege level: `set -privilege advanced`
- c. Move the MDV volumes, one at a time: `volume move start -volume mdv-volume -destination-aggregate aggr-on-new-node -vserver vservers-name`

The following example shows the command and output for moving MDV_CRS_d6b0b313ff5611e9837100a098544e51_A to aggregate data_a3 on node_A_3.

```
cluster_A::> vol move start -volume
MDV_CRS_d6b0b313ff5611e9837100a098544e51_A -destination-aggregate
data_a3 -vserver cluster_A

Warning: You are about to modify the system volume
        "MDV_CRS_d6b0b313ff5611e9837100a098544e51_A". This might
cause severe
        performance or stability problems. Do not proceed unless
directed to
        do so by support. Do you want to proceed? {y|n}: y
[Job 494] Job is queued: Move
"MDV_CRS_d6b0b313ff5611e9837100a098544e51_A" in Vserver "cluster_A"
to aggregate "data_a3". Use the "volume move show -vserver cluster_A
-volume MDV_CRS_d6b0b313ff5611e9837100a098544e51_A" command to view
the status of this operation.
```

- d. Use the volume show command to check that the MDV volume has been successfully moved: `volume show mdv-name`

The following output shows that the MDV volume has been successfully moved.

```
cluster_A::> vol show MDV_CRS_d6b0b313ff5611e9837100a098544e51_B
Vserver      Volume      Aggregate    State      Type      Size
Available Used%
-----
cluster_A    MDV_CRS_d6b0b313ff5611e9837100a098544e51_B
aggr_a2      online     RW          10GB
9.50GB      0%
```

- e. Return to admin mode: `set -privilege admin`

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system- without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.